

AI RAG Eval Platform 0.1.0 OAS 3.1

/openapi.json

default

- POST** /eval/regression Regression Eval
- GET** /health Health
- GET** /stats Stats
- POST** /ingest Ingest
- POST** /query Query
- POST** /query_guarded Query Guarded
- POST** /eval/run Eval Run
- GET** /metrics Metrics

Schemas

Google Cloud AI RAG Eval Platform Search (/) for resources, docs, products, and more

CLOUD SHELL Terminal (ai-rag-eval-prod) x (ai-rag-eval-prod) x (ai-rag-eval-prod) x +

Public API endpoint:
https://ai-rag-eval-api-69725201265.us-central1.run.app

Example request:
curl -X POST https://ai-rag-eval-api-69725201265.us-central1.run.app/eval/run \\\n-H "Content-Type: application/json" \\\n-d '{"prompt":"What is the capital of France?"}'

EOF

azohaib_0150@cloudshell:~/ai-rag-eval-platform (ai-rag-eval-prod)\$ git add README.md

git commit -m "Fix live API endpoint example"

git push

[main ed3beeb] Fix live API endpoint example

1 file changed, 13 insertions(+)

Enumerating objects: 5, done.

Counting objects: 100% (5/5), done.

Delta compression using up to 4 threads

Compressing objects: 100% (3/3), done.

Writing objects: 100% (3/3), 431 bytes | 431.00 KiB/s, done.

Total 3 (delta 2), reused 0 (delta 0), pack-reused 0

remote: Resolving deltas: 100% (2/2), completed with 2 local objects.

To https://github.com/Electricpaper77/ai-rag-eval-platform.git

code8e7f..ed3beeb main -> main

azohaib_0150@cloudshell:~/ai-rag-eval-platform (ai-rag-eval-prod)\$ curl -X POST https://ai-rag-eval-api-69725201265.us-central1.run.app/eval/run \\\n-H "Content-Type: application/json" \\\n-d '{"prompt":"What is the capital of France?"}'

{"status":"ok","results":[{"id":"refund_1","question":"What is the refund policy?","latency_ms":126,"num_citations":0,"top_source":null,"answer":"I cannot answer this question based on the available documents."}, {"id":"ship_1","question":"How long does shipping take?","latency_ms":193,"num_citations":0,"top_source":null,"answer":"I cannot answer this question based on the available documents."}, {"id":"support_1","question":"What are the support hours?","latency_ms":126,"num_citations":0,"top_source":null,"answer":"I cannot answer this question based on the available documents."}], "stats":{"total_questions":3,"questions_with_citations":0,"hit_rate_pct":0.0,"avg_latency_ms":148.3}} azohaib_0150@cloudshell:~/ai-rag-eval-platform (ai-rag-eval-prod)\$


```
https://console.cloud.google.com/welcome?project=ai-rag-eval-prod&cloudshell=true

Google Cloud AI RAG Eval Platform Search (/) for resources, docs, products, and more Search

CLOUD SHELL
Terminal (ai-rag-eval-prod) x (ai-rag-eval-prod) x (ai-rag-eval-prod) x + Open Editor

Gemini CLI is available in Cloud Shell terminal! Type gemini to try it. Learn more Don't show again Dismiss

(status:"ok")azohaib_0150@cloudshell:~/ai-rag-eval-platform (ai-rag-eval-prod)$ ./scripts/run_benchmarks.sh
Running benchmark with 10 prompts...
{
  "prompts": 10,
  "pass_rate": 0.0,
  "p50_latency_ms": 3.75,
  "p95_latency_ms": 5.41,
  "throughput_rps": 253.47
}
[artifact] benchmarks/benchmark_1773304418.json written
Running benchmark with 50 prompts...
{
  "prompts": 10,
  "pass_rate": 0.0,
  "p50_latency_ms": 3.24,
  "p95_latency_ms": 6.22,
  "throughput_rps": 288.58
}
[artifact] benchmarks/benchmark_1773304419.json written
Running benchmark with 100 prompts...
{
  "prompts": 10,
  "pass_rate": 0.0,
  "p50_latency_ms": 3.6,
  "p95_latency_ms": 4.9,
  "throughput_rps": 270.46
}
[artifact] benchmarks/benchmark_1773304419.json written
Running benchmark with 200 prompts...
{
  "prompts": 10,
  "pass_rate": 0.0,
  "p50_latency_ms": 3.85,
  "p95_latency_ms": 4.6,
  "throughput_rps": 254.63
}
[artifact] benchmarks/benchmark_1773304419.json written
azohaib_0150@cloudshell:~/ai-rag-eval-platform (ai-rag-eval-prod)$
```